

How an alert moves through the system

The end-to-end path every alert takes, who owns each stage, and what happens at the three points where a decision is made.

THE RULE THE WHOLE SYSTEM IS BUILT AROUND

The AI classifies. Code decides. The reasoning layer gives a classification, a confidence score and its working. It does not choose what happens to the alert. That is decided by fixed rules in ordinary code, which the AI cannot influence, argue with or talk past.

The pipeline

Every alert goes through all nine stages. Nothing is skipped, and severity never decides whether an alert gets looked at — only how urgently it is escalated.

- 1 Alert arrives**

Picked up the moment it lands. There is no queue to sit in and no analyst choosing what to work on first.

OWNER: **DETECTION TOOL** — EMAIL SECURITY, SIGN-IN MONITORING, EDR OR CLOUD POSTURE
- 2 Parsed and normalised**

The rule, what that rule was written to catch, every person, machine, address and file, and every line of raw evidence. Anything missed here is never enriched later.

OWNER: **THE AGENT**
- 3 Linked to other alerts**

Checks whether this shares infrastructure or people with anything else. Indicators every threat feed calls harmless are excluded, so two reports about the same newsletter are a coincidence rather than a campaign.

OWNER: **THE AGENT**
- 4 Background checks — five sources at once**

Identity, asset inventory, threat intelligence, twelve months of history, and behavioural baselines, all queried in parallel. A source that fails is recorded as a data gap: 20 confidence points if it is silent, 8 if it half-answers, and the result is capped at 90 so the agent can never claim certainty while working blind.

OWNER: **THE AGENT** • NEVER SKIPPED
- 5 Never-close rules — first decision point**

Fifteen situations that always go to a human. Checked in code **before the AI is asked anything**, so the AI cannot argue its way past one. Five of the fifteen depend on EDR data.

OWNER: **DETERMINISTIC CODE** — THE AI HAS NO VOTE
- 6 The AI reasons**

It reads everything gathered so far and returns a fixed structure: one of five classifications, a confidence score, the competing innocent and malicious explanations, and the single fact that would settle which is true. A refusal, a blocked answer or a cut-off answer counts as a failure to triage, not as an all-clear.

OWNER: **THE REASONING MODEL** • TYPICALLY 8–25 SECONDS

7

The policy gate — second decision point

Applies the confidence penalties, then tests whether every close condition holds. Any single failure blocks the close. This is where the outcome is actually set.

OWNER: **DETERMINISTIC CODE** — THE AI HAS NO VOTE

8

Handover package assembled

If the alert is going to a person, sixteen mandatory sections are built and checked for completeness. A missing section is recorded as a defect, not tolerated.

OWNER: **THE AGENT**

9

Recorded, and routed to a person if needed

The decision is written back to FortiSIEM, added to the audit log, and — depending on the outcome — put in front of a Tier 2 analyst or the supervisor. A decision nobody recorded did not happen.

OWNER: **THE AGENT**, THEN A NAMED HUMAN

The three outcomes

There is no fourth. Every alert ends in one of these, and each goes somewhere different.

CLOSE No human needed. Only allowed when every close condition held: a specific innocent explanation, the deciding fact obtained and supporting it, no never-close rule, every source answered, and confidence at or above 85. ----- → Supervisor checks a sample	ESCALATE The default, and not a failure. Any never-close rule, anything suggesting it worked, or a close test that did not fully pass. Escalating too much costs minutes; closing too much costs an incident. ----- → Tier 2 Analyst investigates	HOLD Answerable, but not by the AI right now — usually someone has to be asked. Never a soft close: it records what it is waiting for, who owes the answer, and expires after four hours into an escalation. ----- → Supervisor decides
--	--	--

Who does what after the decision

OUTCOME	WHO PICKS IT UP	WHAT THEY DO	HOW IT ENDS
Escalate	Tier 2 Analyst	Reads the sixteen-section package and investigates. No re-gathering the background.	Stays with Tier 2. Nobody else may move it.
Hold	AI Supervisor	Gets the missing answer, then either closes it or sends it to Tier 2.	Becomes a close or an escalation, recorded in the supervisor's name.
Close — sampled	AI Supervisor	Confirms the AI was right, or marks it wrong and says why.	Confirmed, or reopened as a hold or escalation.
Close — VIP or critical asset	AI Supervisor	Reviewed every time , never sampled, whenever a VIP, a privileged account or a business-critical system is involved.	As above, but mandatory.
Failed safe	Whoever is on duty	The reasoning layer was unreachable. The alert is routed to people and an alarm raised.	Handled manually. Nothing is lost or silently closed.

Who is allowed to do what

Enforced by the server, not just hidden in the interface. A request from the wrong role is refused even if someone bypasses the screen entirely.

ACTION	TIER 2 ANALYST	AI SUPERVISOR	AUTHORITY IF NOT
Read any case and its full package	Yes	Yes	—
See the AI's reasoning in detail	No	Yes	Supervisor duty
Decide a held alert	No	Yes	Supervisor duty
Review a closure the AI made	No	Yes	Supervisor duty

ACTION	TIER 2 ANALYST	AI SUPERVISOR	AUTHORITY IF NOT
Approve a lesson the AI learns from	No	No	Security Administrator
Change autonomy, thresholds or kill switches	No	No	SOC Manager, via change control

The last two rows are deliberate. Separating "who spots a problem" from "who changes the system because of it" is the control — and neither role in the console today holds those rights.

When something breaks

A background source is down. It is written down as a data gap and confidence drops. A clean auto-close typically flips to a hold. The gap is never quietly ignored, and the AI never claims full confidence on incomplete evidence.

The reasoning layer is unreachable. The alert routes to the human queue and an alarm is raised. Retries and fallback models are tried first. Nothing is closed on a guess.

The alert contains instructions aimed at the AI. Attacker-written text is fenced off as evidence, never instruction. The AI is told to treat its presence as a sign of malice, quote it, and reason as if it were not there. Even a fully fooled AI could not close the alert, because the never-close rules were already checked in code.

Describes the system as built. Generated from the current code — the pipeline in `server.js`, the never-close rules and policy gate in `lib/agent.js`. If this and the code disagree, the code is right.