

Your workflow as the AI Supervisor

A rotating duty, not a new job. You supervise the agent instead of doing its work: you decide what it was not confident enough to decide, and you check that what it closed on its own was right to close.

WHY THIS DUTY EXISTS

An autonomous closure is not the end of the story. **Something has to check the checker.** Without this role the agent's accuracy is an assumption; with it, the accuracy is measured, and the measurement is what earns the agent more autonomy later.

The two things that land with you

HELD ALERTS

The agent stopped and asked for a human. It will tell you exactly what it is waiting for and who can answer it. Your job is to get that answer and make the call.

SAMPLED CLOSURES

The agent closed these by itself. Your job is to confirm it was right to. Every closure is reviewed in this build; in production it would be a valid sample of a far larger number.

Two closures are always reviewed, never sampled. Anything involving a VIP, a privileged or service account, or a business-critical system is put in front of you every single time. Blast radius outranks confidence.

Deciding a held alert

- 1 Read why it stopped**

A hold always names what it is waiting for, who owes the answer, and when it expires. If it names neither, that is a defect worth reporting.
- 2 Get the missing answer**

Usually a message to the person involved or the system owner. This is often two minutes of human time on something no amount of telemetry could have settled.
- 3 Close it, or send it to Tier 2**

Those are the only two options. Whichever you pick is recorded in your name, and the incident record is rewritten to show that you decided it, not the agent.

RECORDED AS: **A NAMED HUMAN DECISION**

A hold never quietly expires into a closure. After four hours it escalates to Tier 2 by itself. Time passing is not evidence that nothing was wrong.

Reviewing a closure

Two answers exist. There is no "probably fine".

CORRECT

You agree it was safe to close. Recorded against your name and counted towards the agreement rate the autonomy gate depends on.

WRONG

Marking it wrong actually undoes the closure. You must supply three things before it will save.

YOU MUST SAY	WHY IT IS REQUIRED
What it should have been	Which of the five classifications was correct.
What happens to the alert now	Send it to Tier 2, or put it back on hold for someone to look at. Closing it again is not offered, because closing it is what was wrong.
Why it was wrong	Required, and the box will not submit without it. This sentence is the only part the agent can actually learn from — "it just was" teaches nothing.

What happens to your feedback

- 1

Your reason becomes a candidate lesson

Marked **Pending review**, attributed to you, and automatically checked against existing approved lessons for the same alert type so a contradiction is flagged rather than silently merged.
- 2

A named administrator decides whether it is adopted

Not you. An analyst submits feedback; someone else decides whether the agent learns from it. That separation is the whole control — it is what stops one person quietly retraining the system.
- 3

Once approved, it is used and can be reversed

Approved lessons are given to the agent on every future alert of that type. They are versioned, attributable and reversible. There is no black-box learning anywhere in this system.

Known gap in the current build. No role in the console holds the right to approve a lesson, so candidates accumulate as "Pending review" and are never adopted. Your feedback is still recorded, attributed and visible — but until the Security Administrator role is implemented, nothing you write changes the agent's behaviour. Worth knowing before you rely on it.

What you can and cannot do

YOU CAN ✓ Everything a Tier 2 Analyst can do ✓ See the AI's full reasoning — the "How it decided" tab is yours alone ✓ Decide any held alert: close it or send it up ✓ Mark a closure correct or wrong, and reopen it ✓ Watch the confidence spread and the hold rate	YOU CANNOT ✗ Move an alert already escalated to Tier 2 — an analyst may be working it ✗ Approve your own lesson, or anyone else's ✗ Change autonomy, thresholds or kill switches ✗ Close a held alert without recording a reason
---	--

Your dashboard

NUMBER	WHAT IT TELLS YOU	ACT WHEN
Waiting for you to check	Held alerts plus closures not yet reviewed.	Anything above zero.
Attacks missed	A genuinely malicious alert that got closed. The single most important number here.	Immediately, at any value above zero.
% matched an expert	How often the agent agreed with the recorded human verdict.	Below 95% — the gate needs 95 before the agent is trusted further.
% you marked wrong	How often you are overruling it, out of what you have checked.	Rising over time, which suggests drift.

A disagreement is not automatically a fault. If the agent held an alert because it honestly could not obtain the deciding fact, and you then obtained it by asking someone, the agent behaved correctly and the disagreement is the system working.

Describes the AI Supervisor role as built. Roles and permissions are enforced by the server, so a refused action is refused everywhere, not only on screen.